

week_3 logistic regression

classification and representation

• Introduction

◦ task

- facing discrete data, we usually need to make a **classification**.
- begin with **binary classification problem** ➡ further **multiclass**
- Hence, $y \in \{0,1\}$. 0 is also called the negative class, and 1 the positive class

◦ method

- linear regression doesn't work
- use **logistic regression**

• hypothesis representation

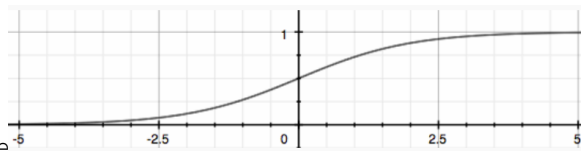
◦ Sigmoid function (Logistic function)

$$h_{\theta}(x) = g(\theta^T x)$$

$$z = \theta^T x$$

$$g(z) = \frac{1}{1 + e^{-z}}$$

■



■ image

- $h_{\theta}(x)$ give us the **probability** that our output is 1. Just interpret it in **Conditional probability**.

$$h_{\theta}(x) = P(y = 1|x; \theta) = 1 - P(y = 0|x; \theta)$$

$$P(y = 0|x; \theta) + P(y = 1|x; \theta) = 1$$

■

- when can we get $h_{\theta}(x) > 0.5$? studying the hypothesis closer.

$$\theta^T x \geq 0 \Rightarrow y = 1$$

$$\theta^T x < 0 \Rightarrow y = 0$$

- we got this:
- and take the term: **Decision Boundary**

◦ Decision Boundary

- the line that separates the area where $y = 0$ and where $y = 1$.
- the property of $h_{\theta}(x)$, once the θ is decided, we got decision boundary. It doesn't come from data set.
- **Image!** remember that decision boundary comes from $\theta^T x$, it can be linear or more complicated based on what polynomial function u is chosen.

Logistic Regression Model

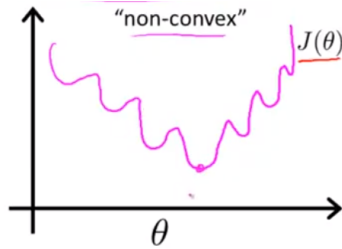
• Cost Function

◦ review:

- **what cost function work.** Cost function is to measure the accuracy of hypothesis. So if our hypothesis comes up with the output that **close to** the actual label, cost function should give value close to 0; while hypothesis gives wrong output, cost function should **penalize** it. That is give high value.
- **how to use cost function to get better θ : Gradient Descent.** When we use GD, we want the cost function *has global optimum* and better *not have local optimum* so that it can converge to the θ we want.

◦ convex:

- non-convex: use the same cost function in linear regression we would get many local optimum. That's not we want. Plot below:



- we use:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

$$\text{Cost}(h_{\theta}(x), y) = -\log(h_{\theta}(x)) \quad \text{if } y = 1$$

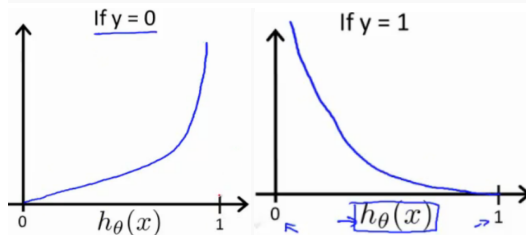
$$\text{Cost}(h_{\theta}(x), y) = -\log(1 - h_{\theta}(x)) \quad \text{if } y = 0$$

- how it work

$$\text{Cost}(h_{\theta}(x), y) = 0 \text{ if } h_{\theta}(x) = y$$

$$\text{Cost}(h_{\theta}(x), y) \rightarrow \infty \text{ if } y = 0 \text{ and } h_{\theta}(x) \rightarrow 1$$

$$\text{Cost}(h_{\theta}(x), y) \rightarrow \infty \text{ if } y = 1 \text{ and } h_{\theta}(x) \rightarrow 0$$



- compress it:

$$h = g(X\theta)$$

$$J(\theta) = \frac{1}{m} \cdot \left(-y^T \log(h) - (1 - y)^T \log(1 - h) \right)$$

• Gradient Descent

Repeat {

$$\theta_j := \theta_j - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

}

-

$$\theta := \theta - \frac{\alpha}{m} X^T (g(X\theta) - \vec{y})$$

- vectored version:
- further: how to make partial derivative?

- link: <https://math.stackexchange.com/questions/477207/derivative-of-cost-function-for-logistic-regression>

The reason is the following. We use the notation

$$\theta x^i := \theta_0 + \theta_1 x_1^i + \dots + \theta_p x_p^i.$$

Then

$$\log h_\theta(x^i) = \log \frac{1}{1 + e^{-\theta x^i}} = -\log(1 + e^{-\theta x^i}),$$

$$\log(1 - h_\theta(x^i)) = \log\left(1 - \frac{1}{1 + e^{-\theta x^i}}\right) = \log(e^{-\theta x^i}) - \log(1 + e^{-\theta x^i}) = -\theta x^i - \log(1 + e^{-\theta x^i}),$$

[this used: $1 = \frac{(1+e^{-\theta x^i})}{(1+e^{-\theta x^i})}$, the 1's in numerator cancel, then we used: $\log(x/y) = \log(x) - \log(y)$]

■

Since our original cost function is the form of:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m y^i \log(h_\theta(x^i)) + (1 - y^i) \log(1 - h_\theta(x^i))$$

Plugging in the two simplified expressions above, we obtain

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \left[-y^i (\log(1 + e^{-\theta x^i})) + (1 - y^i) (-\theta x^i - \log(1 + e^{-\theta x^i})) \right]$$

, which can be simplified to:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \left[y_i \theta x^i - \theta x^i - \log(1 + e^{-\theta x^i}) \right] = -\frac{1}{m} \sum_{i=1}^m \left[y_i \theta x^i - \log(1 + e^{\theta x^i}) \right], \quad (*)$$

where the second equality follows from

$$-\theta x^i - \log(1 + e^{-\theta x^i}) = -\left[\log e^{\theta x^i} + \log(1 + e^{-\theta x^i}) \right] = -\log(1 + e^{\theta x^i}).$$

[we used $\log(x) + \log(y) = \log(xy)$]

All you need now is to compute the partial derivatives of (*) w.r.t. θ_j . As

$$\frac{\partial}{\partial \theta_j} y_i \theta x^i = y_i x_j^i,$$

$$\frac{\partial}{\partial \theta_j} \log(1 + e^{\theta x^i}) = \frac{x_j^i e^{\theta x^i}}{1 + e^{\theta x^i}} = x_j^i h_\theta(x^i),$$

■

the thesis follows.

• Advanced Optimization

- require: Advanced Mathematics. Know how to make **partial derivative**.
- "Conjugate gradient", "BFGS", and "L-BFGS" are more sophisticated, faster ways to optimize θ
- how: use the library.

% write a single function that returns both of these

```
function [jVal, gradient] = costFunction(theta)
    jVal = [...code to compute J(theta)...];
    gradient = [...code to compute derivative of J(theta)...];
end
```

% give to the function "fminunc()" our cost function
% for more details, refer to official document.

• Multi-class Classification: One vs All

- Supposed we have N class. Then train a logistic regression classifier $h_\theta(X)$ for each class to predict the probability that $y = i$. That's say now we have N classifier. Each classifier can predict whether the sample is class(j) or not.
- To make a prediction on a new x, pick the class that **maximizes** $h_\theta(X)$

Regularization

- **over-fitting/under-fitting**

- hypothesis fit training set perfectly but fail to predict / map poorly to the trend of data set.
- solution:

- **Reduce features**

1. manually select features to remain
2. use a model selection algorithm

- **Regularization**

- **modify cost function**

- add regularization parameter λ . It's a **penalty** for theta. It determines how much the costs of our theta parameters are inflated.

$$\min_{\theta} \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2$$

- Proper lambda can make hypothesis smoother so that avoid over-fitting. But
 - too large: penalize too strong that make theta close to 0. Finally we may get a flat line (assumed that θ_0 not been regularized) .
 - too small: help little to fix over-fitting.
- for linear regression:
 - new GD here:

Repeat {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j - \alpha \left[\left(\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \right) + \frac{\lambda}{m} \theta_j \right] \quad j \in \{1, 2, \dots, n\}$$

}

- separate out θ_0 from the rest of the parameters because we do not want to penalize θ_0

$$\theta_j := \theta_j (1 - \alpha \frac{\lambda}{m}) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

- and merge θ :
- new normal equation:

$$\theta = (X^T X + \lambda \cdot L)^{-1} X^T y$$

where $L = \begin{bmatrix} 0 & & & & \\ & 1 & & & \\ & & 1 & & \\ & & & \ddots & \\ & & & & 1 \end{bmatrix}$

- benefit: add the term $\lambda \cdot L$, then $X^T X + \lambda \cdot L$ becomes always **invertible**.
 - when $X_{m \times n}$, has $m < n$, the $X^T X$ would be non-invertible.
 - because:
 - a invertible matrix is full rank.
 - $R(X_{m \times n}) \leq \min(m, n) = m$
 - $X^T X$ is $n \times n$ matrix, which full rank should equal to n .

- Suppose matrix A is of size vxw. Then $\text{rank}(A) \leq \min(v, w)$.
- Also suppose matrices A and B with $\text{rank}(A)=a$ and $\text{rank}(B)=b$. Then $\text{rank}(A*B) = \min(a, b)$.

Remember X is a mxn - matrix, now with $m < n$.

Multiplying out $X^T * X$ is of dimension nxn, but $\text{rank}(X^T * X) \leq m < n$.

Invertibility of a matrix requires full rank, which is by now not given. A square matrix that has full rank is also called **regular**, which is maybe why the name **regularization** arises.

- further explanation:
- also work for **logistic regression**:

- new cost function: (just add θ_j^2 at the end)

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

- by solving partial derivative : new GD

Gradient descent

Repeat {

$$\begin{aligned} \rightarrow \theta_0 &:= \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)} \\ \rightarrow \theta_j &:= \theta_j - \alpha \left[\underbrace{\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}}_{\substack{(j = \cancel{x}, 1, 2, 3, \dots, n) \\ \theta_1, \dots, \theta_n}} + \frac{\lambda}{m} \theta_j \right] \leftarrow \\ \} \end{aligned}$$

$\frac{\partial}{\partial \theta_j} J(\theta)$
 \quad
 $\underline{h_{\theta}(x)} = \frac{1}{1 + e^{-\sigma x}}$

-
- remember to separate θ_0 out.