


# Gradient Descent with Large Data set

- before implementing large data set, try small data set first!
  - verify small data set whether it work or not.
    - by plotting learning curve. when **variance** is **high**, we should try large data set.

## Learning with large datasets



- facing problem: **computational expensive for gradient**

## • Stochastic Gradient Descent

- what the GD we used in the past could be called **batch gradient descent**.
- main idea: **using subset to lower computational cost**.
- **algorithm**
  1. randomly shuffle data set
  2. loop for a subset of examples and iterate to descend gradient by:  
2. For  $i = 1 \dots m$

$$\Theta_j := \Theta_j - \alpha (h_{\Theta}(x^{(i)}) - y^{(i)}) \cdot x_j^{(i)}$$

- attribute:
  - may not converge into global minimum but would be close.
  - though we still have to look through whole data set, **but sometimes before going through, we already get convergent result**. This let us can stop training earlier.

## • Mini-Batch Gradient Descent

### Mini-batch gradient descent

→ Batch gradient descent: Use all  $m$  examples in each iteration

→ Stochastic gradient descent: Use 1 example in each iteration

Mini-batch gradient descent: Use  $b$  examples in each iteration

- 
- just using subset to compute too~
- and here comes a **hyper-parameter**,  $b$  would be hard to choose. Pro.Wang suggested that  $b$  could be 10.

Typical values for  $b$  range from 2-100 or so.



For example, with  $b=10$  and  $m=1000$ :

Repeat:


For  $i = 1, 11, 21, 31, \dots, 991$

$$\theta_j := \theta_j - \alpha \frac{1}{10} \sum_{k=i}^{i+9} (h_{\theta}(x^{(k)}) - y^{(k)}) x_j^{(k)}$$


- advantages: use vectorized implementations to parallel computing.
- tuned learning rate  $\alpha$  / check whether GD is converging
  - facing problem:
    - in small data set we compute  $J(\Theta)$  each iteration and after training, we plot the  $J(\Theta)$  to check
    - however computing  $J(\Theta)$  need to scan through whole data set, when our data set is large. the cost is unacceptable.
  - method:
    - instead of computing  $J(\Theta)$ , we compute  $\text{cost}(\Theta, (x^i, y^i))$ 
      - Stochastic gradient descent:
        - $\text{cost}(\theta, (x^{(i)}, y^{(i)})) = \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$
        - During learning, compute  $\text{cost}(\theta, (x^{(i)}, y^{(i)}))$  before updating  $\theta$  using  $(x^{(i)}, y^{(i)})$ . changeable
        - Every 1000 iterations (say), plot  $\text{cost}(\theta, (x^{(i)}, y^{(i)}))$  averaged over the last 1000 examples processed by algorithm.
      - then plot the average\_cost-iteration to examine.
      - pic may like this:(due to average, there are noise in pic)
        - successfully implement:



- bull shit!



- sometimes:




- 
- we need to try a large iterations like 1000→5000, to see whether the problem caused by noise or divergence.

## • Online Learning

- facing **continuous data stream**. (e.g. shopping online) It's not smart to train the model periodically. Instead, we need to update our model **instantly**.
- **main idea**:
  - **don't save data. Use and discard.**
  - **update  $\theta$  by feeding example.**
- **algorithm**
  - like SGD, just compute GD with one example.
  - for instance, linear regression :
    - $\Theta: \Theta - \alpha(h(X) - Y)X$
    - here y can be user behavior like whether click or not.
- advantages:
  - save place ( data used as one-time)
  - dynamic (sensitive to user preference changing)
- more examples:
  1. improve CTR (Click-Through-Rate)
  2. contents recommendation ( combining with collaborative filter)
  3. Commodity pricing

## • Map and Reduce

- split the task and distribute it
- algorithm is MapReduceable: it can be expressed as **computing sums of functions** over the training set.



- on multiple cores or machines